


Scheduling

Scheduling




Question: How the OS decides which thread should run on cpu next?

Scheduling: deciding which threads are given access to cpu from time to time

Execution Phases of a Process

CPU Burst: the process uses the processor to execute certain instructions before it is no longer ready

I/O Burst: when a process enters the blocked state waiting for I/O




CPU burst

I/O Burst



CPU burst

I/O Burst



time

Scheduling Policies

- Minimize waiting time
 - Processes should not wait long in the ready queue
- Maximize CPU utilization
 - CPU should not be idle
- Maximize throughput
 - Complete as many processes as possible per unit time
- Minimize response time
 - CPU should respond immediately
- Fairness
 - Give each process a fair share of CPU

Scheduling Policies

Waiting time for process p : \sum time before p gets scheduled

Average waiting time: Average of all processes' waiting time.


Completion time (response time): waiting time + run time /
time of completion – time of arrival

Average completion time (response time): Average of all processes' completion time

First-Come, First-Served (FCFS) Scheduling


- First-Come First-Served (FCFS)
 - Process that requests the CPU **FIRST** is allocated the CPU **FIRST**.
 - Also called **FIFO** (First in First Out) “Run until done”
 - In early systems, FCFS meant one program scheduled until done (including I/O)

Process	Arrival Time	Burst Time
P1	0	7
P2	2	4
P3	4	2



First-Come, First-Served (FCFS) Scheduling

Process	Arrival Time	Burst Time
P1	0	7
P2	2	4
P3	4	2




waiting time for P1 =
waiting time for P2 =
waiting time for P3 =
average waiting time =
average completion time =

completion time for P1 =
completion time for P2 =
completion time for P3 =
average completion time =

First-Come, First-Served (FCFS) Scheduling

Process	Arrival Time	Burst Time
P1	0	7
P2	0	4
P3	0	2

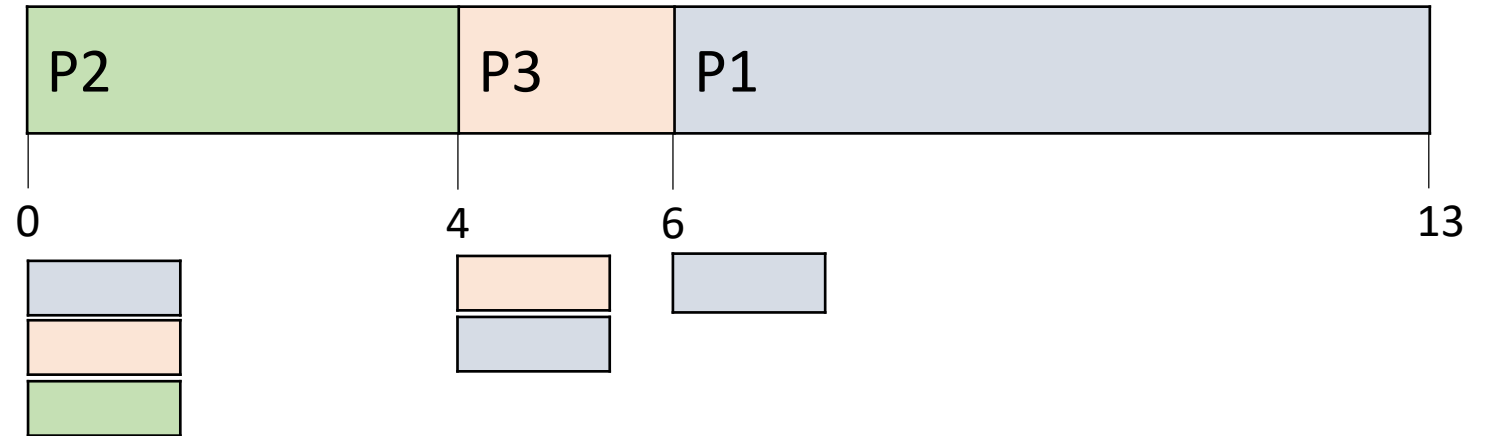


waiting time for P1 =
waiting time for P2 =
waiting time for P3 =
average waiting time =
average completion time =

completion time for P1 =
completion time for P2 =
completion time for P3 =
average completion time =

First-Come, First-Served (FCFS) Scheduling

Process	Arrival Time	Burst Time
P1	0	7
P2	0	4
P3	0	2



waiting time for P1 =
waiting time for P2 =
waiting time for P3 =
average waiting time =
average completion time =

completion time for P1 =
completion time for P2 =
completion time for P3 =
average completion time =

First-Come, First-Served (FCFS) Scheduling

- Pros and Cons

- + simple – Just a FIFO queue

- waiting time depends on arrival order. Bad for short processes

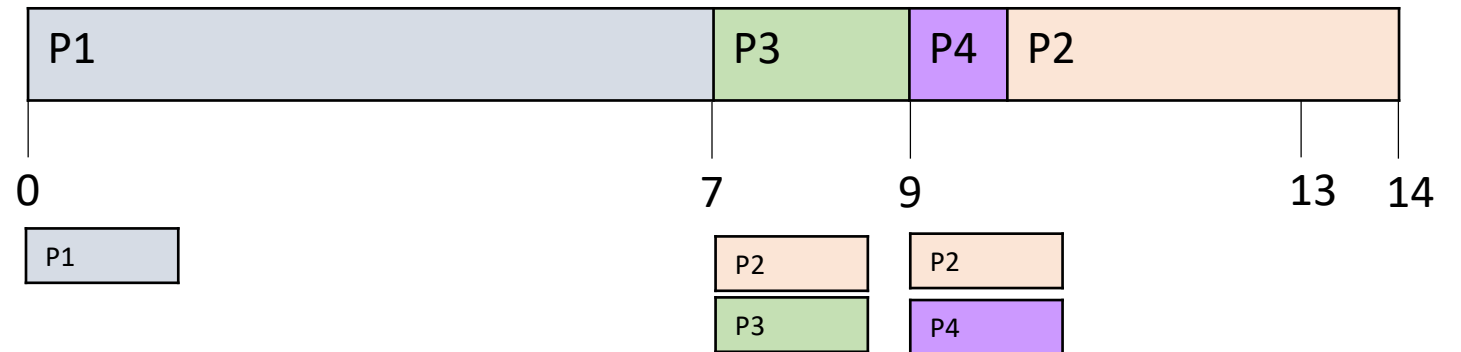
- short processes get stuck waiting for long process to complete (convoy effect)

Shortest Job First (SJF)

- Shortest Job First (SJF)


- schedule the process that has the shortest burst time
 - Randomly pick one if all the processes have the same burst time.
- sometimes called shortest time to completion first
- **No preemption.** Once a process is holding a resource (i.e. once its request has been granted), then that resource cannot be taken away from that process until the process voluntarily releases it.

Process	Arrival Time	Burst Time
P1	0	7
P2	2	4
P3	4	2
P4	8	1



Shortest Job First (SJF)

Process	Arrival Time	Burst Time
P1	0	7
P2	2	4
P3	4	2
P4	8	1



waiting time for P1 =
waiting time for P2 =
waiting time for P3 =
waiting time for P4 =
average waiting time =

completion time for P1 =
completion time for P2 =
completion time for P3 =
completion time for P4 =
average completion time =

Shortest Remaining Time First (SRTF)

- Shortest Remaining Time First (SRTF)/ SJF with preemption
 - If a new process arrives with a shorter burst time than the remaining of the process then schedule the new process

Proce ss	Arrival Time	Burst Time
P1	0	7
P2	2	4
P3	4	2
P4	8	1


Shortest Remaining Time First (SRTF)

Process	Arrival Time	Burst Time
P1	0	7
P2	2	4
P3	4	2
P4	8	1

Process	Arrival Time	Burst Time
P1	0	7
P2	2	4
P3	4	2
P4	8	1




P2 arrives
CPU burst of P2 < remaining time of P1
Hence, preempt P1




Shortest Remaining Time First (SRTF)

Process	Arrival Time	Burst Time
P1	0	7
P2	2	4
P3	4	2
P4	8	1

Process	Arrival Time	Burst Time
P1	0	7
P2	2	4
P3	4	2
P4	8	1



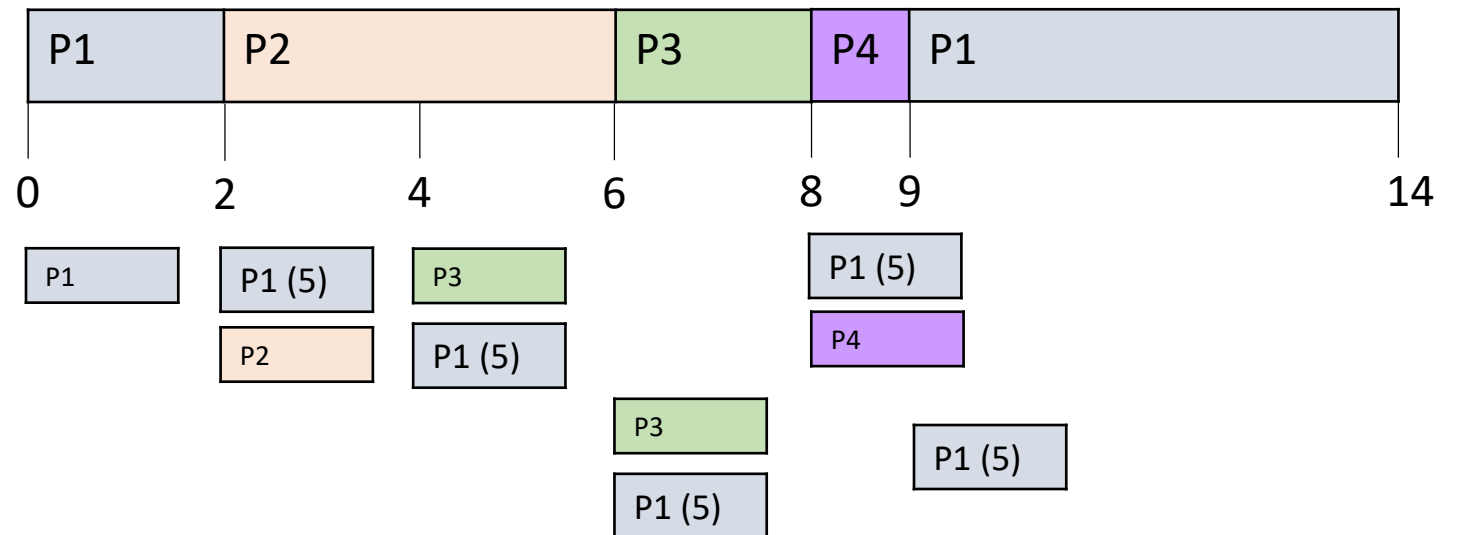
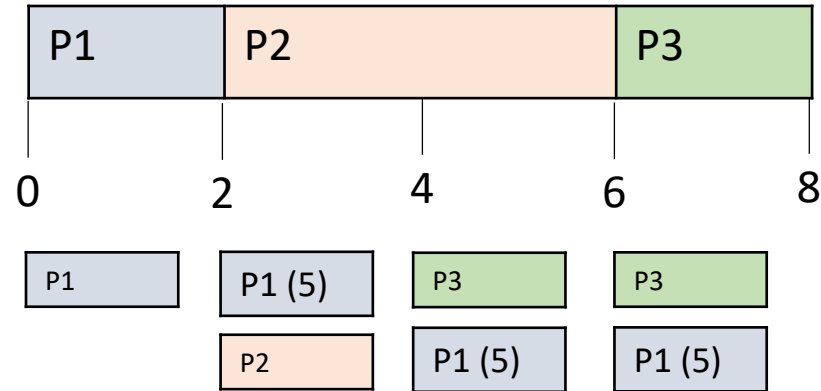
P3 arrives
 CPU burst of P3 = remaining time of P2
 Hence, no preemption



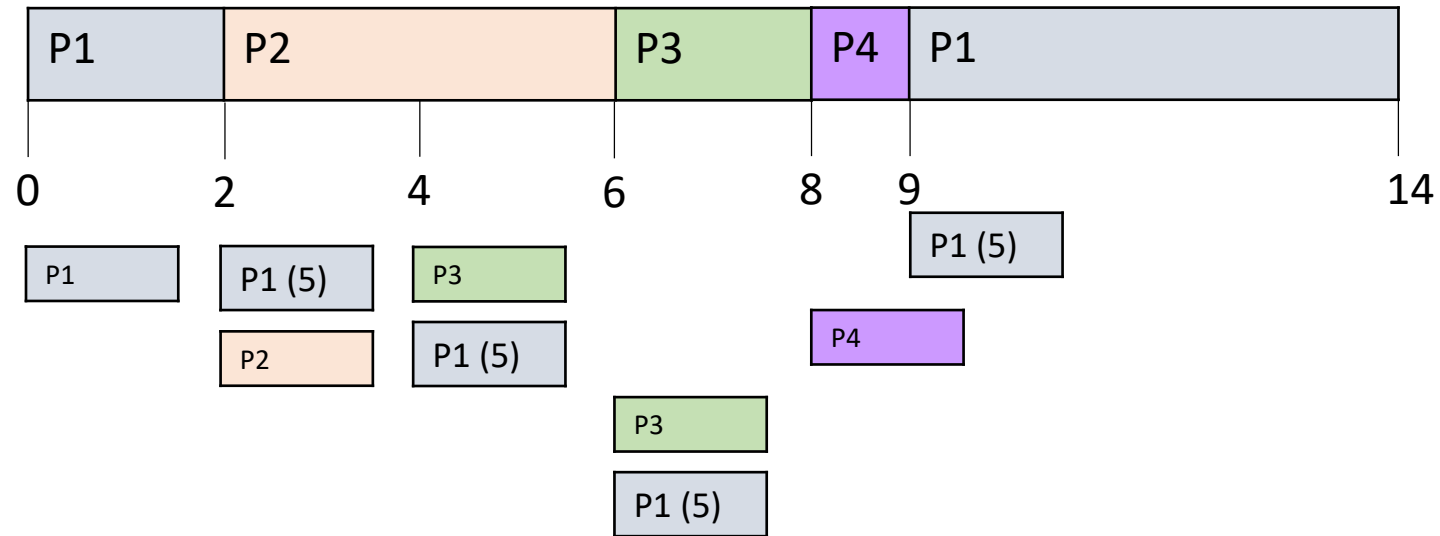
Shortest Remaining Time First (SRTF)

Process	Arrival Time	Burst Time
P1	0	7
P2	2	4
P3	4	2
P4	8	1

Process	Arrival Time	Burst Time
P1	0	7
P2	2	4
P3	4	2
P4	8	1



Shortest Remaining Time First (SRTF)



waiting time for P1 =
waiting time for P2 =
waiting time for P3 =
waiting time for P4 =
average waiting time =

completion time for P1 =
completion time for P2 =
completion time for P3 =
completion time for P4 =
average completion time =

Shortest Remaining Time First (SRTF)

- Pros and Cons

- + Optimal average response time

- you have to predict future

- Unfair

- Multiple small jobs this can lead to starvation

- Large jobs may never get to run

SJF / SRTF

- SJF / SRTF are the best you can do at minimizing average response time
 - Provably optimal (SJF among non-preemptive, SRTF among preemptive)
 - SRTF is always at least as good as SJF
- Comparison of SRTF with FCFS
 - What if all jobs the same length?
 - SRTF becomes the same as FCFS (i.e. FCFS is best can do if all jobs the same length)
 - What if jobs have varying length?
 - SRTF: short jobs not stuck behind long ones

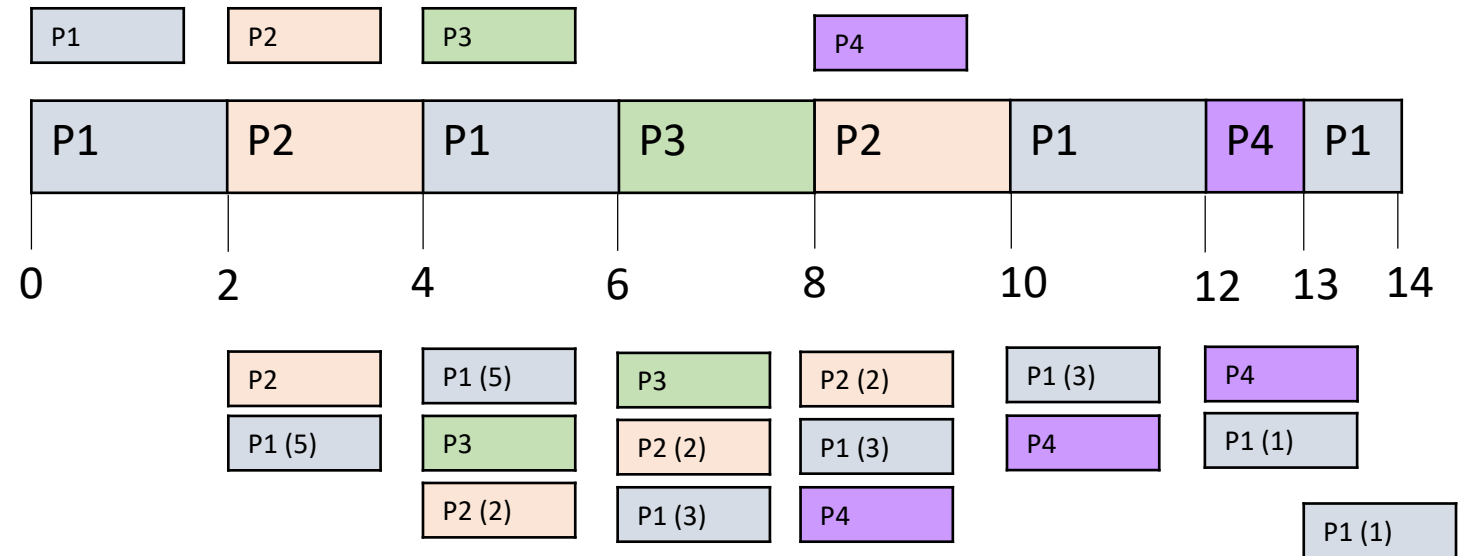
Robin Round Scheduling

- Robin Round Scheduling
 - Each process gets a small unit of CPU time (time quantum / time slice)
 - After quantum / slice expires the process is moved to a FIFO

Proce ss	Arrival Time	Burst Time
P1	0	7
P2	2	4
P3	4	2
P4	8	1

Robin Round Scheduling

Process	Arrival Time	Burst Time
P1	0	7
P2	2	4
P3	4	2
P4	8	1



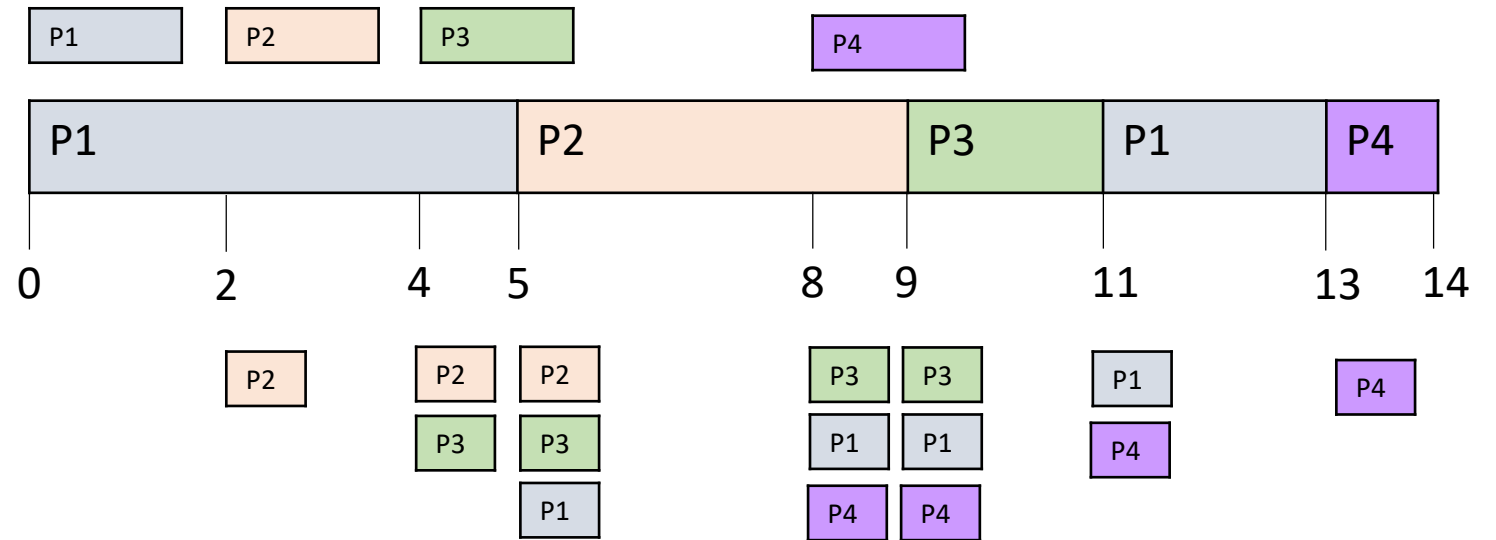
average waiting time =
average response time =
Context switches =

Time slice / quantum (q)= 2

Robin Round Scheduling

Process	Arrival Time	Burst Time
P1	0	7
P2	2	4
P3	4	2
P4	8	1

average waiting time =
average response time =
Context switches =



Time slice / quantum (q) = 5

Robin Round Scheduling

P1: Burst Length 8 (Arrival: 0)


P2: Burst Length 2 (Arrival: 0)

- $q = 8$



Average Response time = $(8 + 10) / 2 = 9$

- $q = 4$




Average Response time = $(10 + 6) / 2 = 6$

Robin Round Scheduling

P1: Burst Length 2 (Arrival: 0)


P2: Burst Length 2 (Arrival: 0)

- $q = 8$



$$\text{Average Response time} = (2 + 4) / 2 = 3$$

- $q = 4$




$$\text{Average Response time} = (2 + 4) / 2 = 3$$

Robin Round Scheduling

P1: Burst Length 2 (Arrival: 0)


P2: Burst Length 2 (Arrival: 0)

- $q = 2$



$$\text{Average Response time} = (2 + 4) / 2 = 3$$

- $q = 1$



$$\text{Average Response time} = (3 + 4) / 2 = 3.5$$



Robin Round Scheduling

- How to choose q ?
 - Too Large: FCFS
 - Too Small: High Context Switches – Increased Overhead
 - q must be large with respect to the number of context switch

Robin Round Scheduling

Pros and Cons:

- + Fair (Each process gets a fair chance to run on the CPU)
- + Faster response time
- Increased context switching (increased overhead)
- Bad when the processes have same burst length